```
~/Desktop/cw1-Build/bin $ ./lrgLinReg -n TestData1.txt
Using normal equation yields the following results:
theta0 = 3.01749
theta1 = 1.9966
~/Desktop/cw1-Build/bin $ ./lrgLinReg -n TestData2.txt
Using normal equation yields the following results:
theta0 = 2.01749
theta1 = 2.9966
~/Desktop/cw1-Build/bin $ ./lrgLinReg -g TestData1.txt 1000000 0.001 0.0 0.0
Using gradient descent yields the following results:
theta0 = 3.01749
theta1 = 1.9966
~/Desktop/cw1-Build/bin $ ./lrgLinReg -g TestData2.txt 1000000 0.001 0.0 0.0
Using gradient descent yields the following results:
theta0 = 2.01749
theta1 = 2.9966
```

Above is a screenshot of results when running my command line tool. As you can see the theta values are very close to the true values. For the gradient descent I used:

epochs = 1,000,000
learning_rate = 0.001
theta0_guess = 0.0
theta1_guess = 0.0